AI Agent Architecture Document

Project Title: Personalized Email Generation Assistant

Author: ABHIJEET KUMAR

Enrollment: 23324001

Date: 03-11-2025

1. Overview

The **AI Email Assistant** is a modular system designed to automate the creation of formal and academic outreach emails using a **fine-tuned LLaMA-3.2-3B-Instruct model**.
It integrates structured prompting, reasoning-based text generation, and quantitative evaluation into a single workflow.

The agent is capable of **reasoning**, **planning**, and **executing** tasks related to text composition based on structured input fields provided by the user.

2. Core Components

Component	Description	File/Module
Prompt Construction Module	Converts structured user fields (professor name, subject, bio, and goal) into a natural-language instruction-style prompt.	generate_email.py
LLM Adapter Module	Loads base LLaMA-3 model, attaches LoRA fine-tuned adapter, and handles inference.	llm_adapter.py
Evaluation Module	Computes BLEU, ROUGE-L, BERTScore, and semantic similarity for generated vs. reference outputs.	evaluate_generation.py
Interactive CLI Interface	Provides user interface for generating and saving emails or running evaluation mode.	email_assistant2.py
Fine-Tuning Script	Handles LoRA-based model adaptation with quantization for efficiency.	train_lora.py

3. Interaction Flow

Step-by-Step Pipeline

1. **User Input:**
   The user provides structured data — professor's name, subject, short bio, goal, and tone — through the command-line interface.
2. **Prompt Generation:**
   The system constructs a detailed instruction prompt in the following format:
3. ### Instruction:
4. Write a formal outreach email using:
5. prof_name: ...
6. subject: ...
7. student_bio: ...
8. goal: ...
9. **Model Inference:**
   - The **LLaMA-3.2-3B-Instruct** model is loaded in quantized (4-bit) mode.
   - **LoRA adapters** fine-tuned on a small custom dataset are attached.
   - The model generates the email using top-p sampling and temperature-based decoding.
10. **Response Extraction:**
    The model's output is parsed to isolate the response text following the "### Response:" token.
11. **Email Saving:**
    The generated email and metadata are saved in a timestamped .txt file under generated_emails/.
12. **Evaluation (Optional):**
    The user can evaluate model quality by comparing generated outputs with ground-truth references via semantic and text-similarity metrics.

4. Model Architecture

Base Model

- **Model Name:** meta-llama/Llama-3.2-3B-Instruct
- **Architecture:** Transformer-based autoregressive language model
- **Parameters:** 3 Billion
- **Instruction-Tuned:** Yes
- **Reason for Choice:**
    - Compact yet powerful for natural-language generation
    - Openly available and easily fine-tuned
    - Balanced trade-off between accuracy and computational cost

Fine-Tuning Method

- **Technique:** LoRA (Low-Rank Adaptation)
- **Quantization:** 4-bit (using BitsAndBytes)
- **Training Epochs:** 1
- **Learning Rate:** 2e-4
- **Final Loss:** 1.23
- **Rationale:**
    - LoRA offers efficient adaptation without full retraining
    - Enables use of consumer hardware (single GPU)
    - Achieves high contextual alignment for limited data

5. Data Flow Diagram

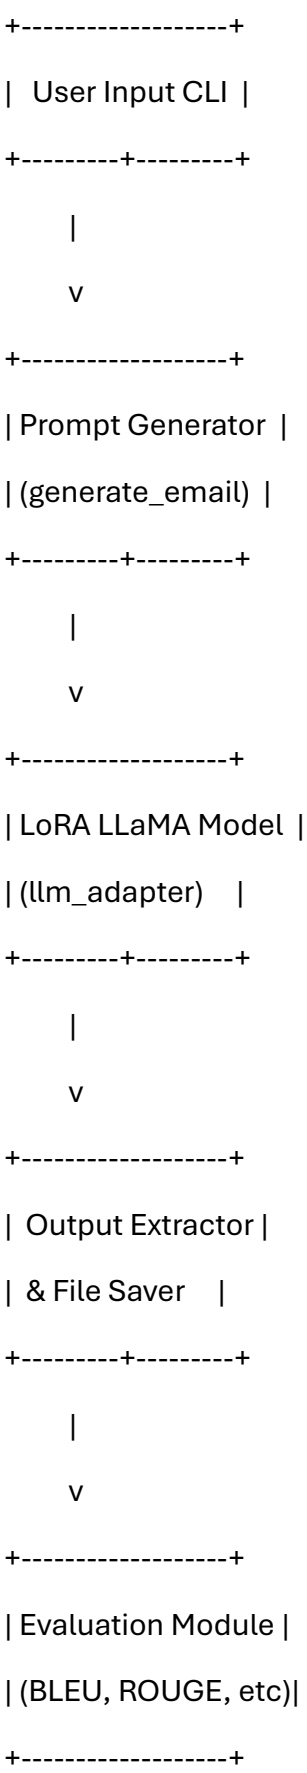

6. Reasoning Behind Design Choices

Design Decision	Justification
LLaMA-3.2-3B as Base Model	Provides instruction-following ability with moderate compute requirements.
LoRA Fine-Tuning	Enables lightweight adaptation and domain specialization.
4-bit Quantization	Reduces VRAM usage and accelerates fine-tuning.
CLI Interface	Simplifies interaction during prototype phase.
Evaluation Metrics Suite	Ensures both semantic and syntactic quality assessment.

7. Summary

The **AI Email Assistant** architecture reflects an efficient and modular design suitable for resource-constrained environments.
It demonstrates how **parameter-efficient fine-tuning** and **structured prompting** can create a robust, task-specific AI agent capable of performing reasoning-based automation tasks such as email composition.